

# Blockchain.com SRE Proficiency Test

Times are suggested totals for each stage: they are just indicative of the level of detail we are expecting.

Feel free to use any tools that would be appropriate to delivering a production-ready system.

## Software Development

(~90 minutes)

You are tasked with starting a new project at blockchain.com. You have received the following list of requirements from our product team:

We require a non-critical API service (~90% SLO) for some already-existing internal services to use. The service should accept commands on a TCP port composed of ASCII strings. The server outputs an ASCII string as response.

The following commands are accepted:

- `WHO`: outputs the total number of clients connected;
- `WHERE`: outputs the id of the server (a unique identifier);
- `WHY`: output the string `42`;

Please leave enough time to provide a list of items that would need to be added to the project before it would be considered production-ready. (Add this in a `README.md` next to the source code)

Consider

- Maintainability (code hygiene, code robustness, documentation)
- Deployability
- Observability

In this part we are interested in

- a) code correctness
- b) how you structure a coding project
- c) tooling you add to achieve the above goals
- d) how straightforward it would be for someone else to continue development where you leave off

- e) any decisions you've had to make along the way (please include these in the `README.md`)

In this part, we are not interested in

- a) how it fits into the larger ecosystem
- b) an exhaustive implementation of all the tooling

## System Design

(~30 minutes)

The server you have written in part 1 has now been used in production as part of our non-customer facing micro service architecture for a few months. The engineering team would now like to leverage it as part of our customer-facing infrastructure and as such the target SLO is now 99.95%.

Design a distributed system composed of at least 3 servers. Clients can connect to any server available.

Server response will change to reflect the distributed architecture:

- **WHO:** the outputs should show the total number of clients connected to all available servers
- **WHERE:** the id of the server needs to be a globally unique identifier

Describe the distributed architecture, in particular how you would achieve state replication and fault tolerance

You should justify your ideas in light of the CAP theorem - do you need more information from the engineering / product teams?

**Note:** code is not required. Design the architecture, describe the architecture, use pseudo code if necessary

### Points for discussion afterwards

What considerations would you have if the SLO were 99.9999%?

What changes would you encourage the product team to make in their requirements?

How does this service fit into a larger microservice ecosystem? What tooling will it require to make it useful?